


Upper Confidence Bound Algorithm



Thompson Sampling Algorithm Intuition

The Multi-Armed Bandit Problem



D1



D2



D3



D4



D5

The Multi-Armed Bandit Problem

- We have d arms. For example, arms are ads that we display to users each time they connect to a web page.
- Each time a user connects to this web page, that makes a round.
- At each round n , we choose one ad to display to the user.
- At each round n , ad i gives reward $r_i(n) \in \{0, 1\}$: $r_i(n) = 1$ if the user clicked on the ad i , 0 if the user didn't.
- Our goal is to maximize the total reward we get over many rounds.

Bayesian Inference

- Ad i gets rewards \mathbf{y} from Bernoulli distribution $p(\mathbf{y}|\theta_i) \sim \mathcal{B}(\theta_i)$.
- θ_i is unknown but we set its uncertainty by assuming it has a uniform distribution $p(\theta_i) \sim \mathcal{U}([0, 1])$, which is the prior distribution.
- Bayes Rule: we approach θ_i by the posterior distribution

$$\underbrace{p(\theta_i|\mathbf{y})}_{\text{posterior distribution}} = \frac{p(\mathbf{y}|\theta_i)p(\theta_i)}{\int p(\mathbf{y}|\theta_i)p(\theta_i)d\theta_i} \propto \underbrace{p(\mathbf{y}|\theta_i)}_{\text{likelihood function}} \times \underbrace{p(\theta_i)}_{\text{prior distribution}}$$

- We get $p(\theta_i|\mathbf{y}) \sim \beta(\text{number of successes} + 1, \text{number of failures} + 1)$
- At each round n we take a random draw $\theta_i(n)$ from this posterior distribution $p(\theta_i|\mathbf{y})$, for each ad i .
- At each round n we select the ad i that has the highest $\theta_i(n)$.

Thompson Sampling Algorithm

Step 1. At each round n , we consider two numbers for each ad i :


- $N_i^1(n)$ - the number of times the ad i got reward 1 up to round n ,
- $N_i^0(n)$ - the number of times the ad i got reward 0 up to round n .

Step 2. For each ad i , we take a random draw from the distribution below:


$$\theta_i(n) = \beta(N_i^1(n) + 1, N_i^0(n) + 1)$$

Step 3. We select the ad that has the highest $\theta_i(n)$.


Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm




Thompson Sampling Algorithm



Thompson Sampling Algorithm




Thompson Sampling Algorithm



Thompson Sampling Algorithm

And so on...


Thompson Sampling Algorithm




UCB vs Thompson Sampling

Thompson Sampling Algorithm

UCB



Thompson Sampling



- Deterministic
- Requires update at every round

- Probabilistic
- Can accommodate delayed feedback
- Better empirical evidence

Natural Language Processing (NLP)

Plan of Attack

Plan of Attack

Here's what we will learn:


- Types of Natural Language Processing
- Classical vs Deep Learning Models
- End-to-end Deep Learning Models
- Bag-Of-Words

- Note: Seq2Seq and Chatbots are outside the scope of this course

Types of NLP

Types of NLP

Natural
Language
Processing




Deep
Learning

Classical vs Deep Learning Models

Classical vs Deep Learning Models

Natural
Language
Processing



Deep
Learning

Classical vs Deep Learning Models

Some examples:

1. If / Else Rules (Chatbot)
2. Audio frequency components analysis (Speech Recognition)
3. Bag-of-words model (Classification)
4. CNN for text Recognition (Classification)








Image Source: www.wildml.com







End-to-end Deep Learning Models

End-to-end Deep Learning Models




End-to-end Deep Learning Models






End-to-end Deep Learning Models



End-to-end Deep Learning Models



End-to-end Deep Learning Models



End-to-end
Deep Learning
Models

Bag-Of-Words



...

Catch up?

Inbox

Business



Kirill Eremenko

to me

6:18 pm ...

Hello Kirill,

Checking if you are back to Oz. Let
me know if you are around and keen
to sync on how things are going. I
defo could use some of your
creative thinking to help with mine :)

Cheers,

V

...

Yes, I'm
around.

I'm back!

Sorry, I'm not.



Reply




Forward

Bag-Of-Words

Yes / No

Bag-Of-Words



171,476 words

The Second Edition of the 20-volume Oxford English Dictionary contains full entries for **171,476 words** in current use, and **47,156** obsolete words. To this may be added around **9,500** derivative words included as subentries.



How many words are there in the English language?

<https://en.oxforddictionaries.com/.../how-many-words-are-there-in-the-english-language>

[About this result](#) [Feedback](#)

People also ask

How many words in the English language does the average person know?

Most adult native test-takers range from **20,000–35,000 words**. Average native test-takers of age 8 already know **10,000 words**. Average native test-takers of age 4 already know **5,000 words**. Adult native test-takers learn almost 1 new word a day until middle age. May 29, 2013

Lexical facts - The Economist

<https://www.economist.com/blogs/johnson/2013/05/vocabulary-size>



We have seen that the Oxford English Dictionary contains **171,476 words** in current use, whereas a vocabulary of just **3000 words** provides coverage for around 95% of common texts. If you do the math, that's **1.75% of the total number of words in use!** Mar 14, 2013

How many words in the english language ? How many do i need to ...

<https://www.lingholic.com/how-many-words-do-i-need-to-know-the-955-rule-in-langua...>

Bag-Of-Words

20,000 elements long

SOS
EOS

Special Words

Bag-Of-Words


Hello Kirill, Checking if you are back to Oz. Let me know if you are around ... Cheers, V

← →

20,000 elements long


Bag-Of-Words

Hello Kirill, Checking if you are back to Oz. Let me know if you are around ... Cheers, V




Bag-Of-Words

Hello Kirill, Checking if you are back to Oz. Let me know if you are around ... Cheers, V



Bag-Of-Words

Hello Kirill, Checking if you are back to Oz. Let me know if you are around ... Cheers, V




Training Data:

Hey mate, have you read about Hinton's capsule networks?	→	No
Did you like that recipe I sent you last week?	→	Yes
Hi Kirill, are you coming to dinner tonight?	→	Yes
Dear Kirill, would you like to service your car with us again?	→	No
Are you coming to Australia in December?	→	Yes
...	→	...

Bag-Of-Words

Hello Kirill, Checking if you are back to Oz. Let me know if you are around ... Cheers, V




Training Data:

[1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, ..., 2]	→	No
[1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 2, 0, 0, 0, 1, 0, 0, 1, 0, 0, ..., 0]	→	Yes
[1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, ..., 1]	→	Yes
[1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, ..., 1]	→	No
[1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, ..., 1]	→	Yes
...	→	...

Bag-Of-Words

Hello Kirill, Checking if you are back to Oz. Let me know if you are around ... Cheers, V



[1, 1, 0, 0, 1, 0, 2, 0, 1, 0, 0, 0, 0, 1, 2, 0, 0, 0, 1, 0, 0, 1, 0, 0, ... , 3] → Yes / No ?

20,000 elements long

Training Data:

[1, 1, 0, 0, 0]
[1, 1, 0, 0, 0]
[1, 1, 0, 0, 0]
[1, 1, 0, 0, 0]
[1, 1, 0, 0, 0]





Image Source: www.helloacm.com



No
Yes
Yes
No
Yes
...



Artificial Neural Network Intuition

A Clip From The Today Show, 1994

What is Deep Learning?


What is Deep Learning?



What is Deep Learning?



What is Deep Learning?



Source: mkomo.com


What is Deep Learning?



What is Deep Learning?

STORAGE LIMITS

Estimates based on bacterial genetics suggest that digital DNA could one day rival or exceed today's storage technology.



Source: nature.com

What is Deep Learning?

1 The accelerating pace of change ...




2 ... and exponential growth in computing power ...

Computer technology, shown here climbing dramatically by powers of 10, is now progressing more each hour than it did in its entire first 90 years


COMPUTER RANKINGS

By calculations per second per \$1,000




3 ... will lead to the Singularity


Surpasses brainpower equivalent to that of all human brains combined



Source: Time Magazine



What is Deep Learning?



Geoffrey Hinton

What is Deep Learning?









Image Source: www.austincc.edu




What is Deep Learning?



What is Deep Learning?



What is Deep Learning?



Supervised vs Unsupervised

Supervised vs Unsupervised

Supervised	Artificial Neural Networks	Used for Regression & Classification
	Convolutional Neural Networks	Used for Computer Vision
	Recurrent Neural Networks	Used for Time Series Analysis
Unsupervised	Self-Organizing Maps	Used for Feature Detection
	Deep Boltzmann Machines	Used for Recommendation Systems
	AutoEncoders	Used for Recommendation Systems

Plan of Attack

Plan of Attack

What we will learn in this section:

- The Neuron
- The Activation Function
- How do Neural Networks work? (example)
- How do Neural Networks learn?
- Gradient Descent
- Stochastic Gradient Descent
- Backpropagation

The Neuron

The Neuron





Image Source: www.austincc.edu

The Neuron




Image Source: Wikipedia

The Neuron




Image Source: Wikipedia

The Neuron






Image Source: Wikipedia


The Neuron




The Neuron




The Neuron




The Neuron



The Neuron



The Neuron



The Neuron


Additional Reading:

Efficient BackProp


By Yann LeCun et al. (1998)

Link:


<http://yann.lecun.com/exdb/publis/pdf/lecun-98b.pdf>




The Neuron




The Neuron




The Neuron




The Neuron




The Neuron




The Neuron



The Neuron




The Neuron




The Activation Function


The Activation Function




The Activation Function




The Activation Function




The Activation Function



The Activation Function



The Activation Function




The Activation Function

Additional Reading:

*Deep sparse rectifier
neural networks*


By Xavier Glorot et al. (2011)




Link:

<http://jmlr.org/proceedings/papers/v15/glorot11a/glorot11a.pdf>

The Activation Function



The Activation Function



How do NNs Work?

